

YOONBO CHO

AI Engineer · Applied LLM Systems

Elkins Park, PA · +1 718 541 6008 · yoonbocho1@gmail.com · yoonbo1.github.io · github.com/yoonbo1 · Open to relocation

SUMMARY

AI engineer who ships production LLM systems end-to-end and then operates them: on-device retrieval, eval harnesses with real ship bars, multi-provider backends, and agent safety. Shipped AI features for a sports/venue platform serving 1.2M monthly active users, and built two live AI products solo. 4+ years of senior mobile and full-stack engineering.

WORK EXPERIENCE

Senior Developer — A View From My Seat

Nov 2024 – Aug 2026 · Philadelphia, PA

Sports/venue platform · 1.2M monthly active users · owned native iOS + Android across 80+ releases

- Shipped **vision-based moderation** for user-submitted seat-view photos, auto-flagging or clearing each image for publication, and backfilled structured metadata (view obstructions, for example) across the existing photo corpus.
- Built a **venue-scouting agent** that runs daily, diffs live venue data against stored snapshots, and alerts only on real changes, replacing manual monitoring.

Founder & AI Engineer — Yoonibo LLC

Jul 2022 – Present

Everblue — passive AI journaling iOS app · everblue.me · live on the App Store, launched Jun 2026

- Built an **LLM eval harness** that runs the real production pipeline on synthetic day-payloads and grades against the production rubric (voice lint, specificity, hallucination). Ship bars of 92%+ generation quality and a 100% no-fabrication gate block every prompt change.
- Engineered an **on-device hybrid-retrieval RAG pipeline**: MiniLM ONNX embeddings, BM25, exponential recency decay, and MMR re-ranking, grounding an AI companion in months of personal context. Deliberately no vector database: the per-user corpus doesn't warrant one.
- Built a **multi-provider LLM backend** (Django) routing Gemini and Claude by feature and subscription tier, with retry-and-fallback orchestration, prompt caching, PII redaction on every request, and measured cost cuts of 3–23% per prompt.
- Shipped fully on-device photo captioning (~260MB FP16 **Florence-2** Core ML, custom Expo config plugin, Swift bridge) and debugged its jetsam crashes on 4GB iPhones; then reversed the design when the ~1.1GB bundle stopped earning its cost, replacing it with a gated backend service.

ticketap — Yoonibo LLC

2026

Taste-driven going-out platform · ticketap.com · live

- Built **Tap**, the booking agent, on a hard safety invariant: purchases require an active mandate enforced server-side in the tool executor, never the prompt. A 25-case adversarial suite passes at 100% and gates CI.

SKILLS

AI / LLM: Retrieval-Augmented Generation (RAG), hybrid retrieval (BM25 + dense embeddings, MMR), LLM evaluation harnesses & guardrails, agent safety, Model Context Protocol (MCP), prompt caching & cost control, on-device inference (Core ML, ONNX Runtime)

LLM platforms: Anthropic Claude, Google Gemini, OpenAI

Languages: Python, TypeScript, Swift, Kotlin, SQL

Backend & Mobile: Django, REST & GraphQL APIs, PostgreSQL, Docker, AWS, Heroku, GitHub Actions CI/CD, React Native / Expo, native iOS & Android

EDUCATION

B.S. Computer Science — Rowan University

2019 – 2021

CERTIFICATIONS

TensorFlow Developer Certificate — image classification, NLP, and time-series forecasting