

YOONBO CHO

AI Engineer · Applied LLM Systems

Elkins Park, PA · +1 718 541 6008 · yoonbocho1@gmail.com · yoonbo1.github.io · github.com/yoonbo1 · Open to relocation

SUMMARY

AI engineer who ships production LLM systems end-to-end and then operates them. Ship AI features — vision moderation, agentic tooling — on a sports/venue platform serving 1.2M monthly active users, and built Everblue solo: a live App Store product with an on-device hybrid-retrieval RAG pipeline, an eval harness that gates every prompt and model change, and a privacy-first multi-provider LLM backend. 4+ years of senior mobile and full-stack engineering; three consumer apps solo-shipped to the App Store, two of them LLM-powered.

WORK EXPERIENCE

Senior Developer — A View From My Seat

Nov 2024 – Aug 2026 · Philadelphia, PA

Sports/venue platform · 1.2M monthly active users · native iOS + Android

- Shipped **vision-based moderation** for user-submitted seat-view photos — auto-flagging or clearing each image for publication — and backfilled structured metadata across the existing corpus, including view-obstruction detection.
- Built an **MCP-driven venue-scouting agent** that runs daily on cron, diffs live venue data against stored snapshots, and alerts only on real changes, replacing manual monitoring.
- Built an agentic dev loop that pulls crash reports, drafts candidate fixes for review, and automates test runs.
- Own native iOS (Swift) and Android (Kotlin) development end-to-end — spec, architecture, implementation, QA, release — shipping 80+ updates at feature parity across platforms.
- Automated the seating-chart ingestion pipeline, replacing a manual upload/publish process; a teammate reported it saved a full day of work per chart.

Founder & AI Engineer — Yoonibo LLC

Jul 2022 – Present

Everblue — passive AI journaling iOS app · everblue.me · live on the App Store, launched Jun 2026

- Built an **LLM eval harness** that runs the real production pipeline on synthetic day-payloads and auto-grades against the production rubric (voice lint, specificity, hallucination) plus a recall-fabrication suite. Ship bars of 92%+ generation quality and a 100% no-fabrication gate block every prompt change; passing runs export a distillation JSONL for fine-tuning.
- Engineered an **on-device hybrid-retrieval RAG pipeline** — all-MiniLM-L6-v2 ONNX embeddings + BM25 + cosine similarity, with exponential recency decay and MMR re-ranking — grounding an AI companion in months of personal context. No vector database and no LLM framework: the per-user corpus and linear call graph didn't warrant either, and I can defend the point at which that would change.
- Built a **multi-provider LLM backend** (Django/Heroku) routing Google Gemini and Anthropic Claude Sonnet/Haiku by feature and subscription tier, with retry-and-fallback orchestration, ephemeral prompt caching, and PII redaction before every prompt. Measured cost work — thinking-token caps and cache-stable prompts — cut 3–23% per prompt.
- Shipped **fully on-device photo captioning**: ~260MB FP16 Florence-2 CoreML injected through a custom Expo config plugin and a Swift/Obj-C bridge. Debugged jetsam/OOM crashes on 4GB iPhones down to device-aware Core ML compute units and fixed-shape loading — then reversed the design when honest accounting showed the ~1.1GB bundle wasn't earning its cost, migrating to a gated backend service with EXIF stripping.
- Designed the LLM product behavior: a deterministic voice-lint gate inside the generation retry loop, date-aware memory recall, a source-grounding pass that can only remove ungrounded claims, and honest refusal on blank days rather than fabricated entries.
- Built an autonomous distribution stack — read-only browser-agent scouts (Playwright MCP, parallel per-profile Chrome, one batch-served local LLM) feeding a daily Claude agent that drafts posts behind a human-approved admin API, rewriting its own search queries from signal quality each run.

- Built **Tap**, the booking agent, around a hard safety invariant: purchase execution requires an active mandate enforced **server-side in the tool executor**, never the prompt — a 25-case adversarial mandate suite passes 100% and gates CI.
- Engineered 30-case agentic tool-calling eval suites with LLM-judge scoring, and measured judge variance before trusting verdicts.

Earlier products — Yoonibo LLC2022 – 2024

- **Jibbr** (2024) — OpenAI-powered conversation coach generating interest-tailored prompts and actionable soft-skills feedback.
- **+Luv** (2022) — cross-platform couples app on the iOS App Store and Google Play, with WebSocket real-time chat, OpenAI-generated prompts, photo sharing, and subscriptions.

Software Developer — Fabspeed MotorsportJan 2022 – Jul 2022 · Philadelphia, PA

- Raised site conversion through heatmap-driven UX optimization, launched wholesale.fabspeed.com, and built Python scrapers extracting Google and Yelp reviews at scale.

SKILLS

AI / LLM: Retrieval-Augmented Generation (RAG), hybrid retrieval (BM25 + dense embeddings, MMR re-ranking), LLM evaluation harnesses & guardrails, agentic systems, Model Context Protocol (MCP)

LLM platforms: Anthropic Claude, Google Gemini, OpenAI, prompt caching & cost control, on-device inference (Core ML, ONNX Runtime), embeddings (MiniLM)

Languages: Python, TypeScript, Swift, Kotlin, SQL

Backend & Infra: Django, REST & GraphQL APIs, PostgreSQL, Docker, Amazon Web Services (AWS), Heroku, Continuous Integration/Continuous Deployment (CI/CD) via GitHub Actions

Mobile: React Native / Expo, native iOS & Android, native modules, App Store shipping

EDUCATION

B.S. Computer Science — Rowan University2019 – 2021

CERTIFICATIONS

TensorFlow Developer Certificate — image classification, NLP, and time-series forecasting